

The logo for Oracle Academy. The word "ORACLE" is in a bold, orange, sans-serif font. Below it, the word "Academy" is in a smaller, dark gray, sans-serif font. The entire logo is centered on a light gray background, which is framed by dark gray horizontal bars at the top and bottom.

# ORACLE

## Academy

# Java Programming

9-1

Java Bytecode

**ORACLE**  
Academy



Copyright © 2020, Oracle and/or its affiliates. All rights reserved.

# Objectives

- This lesson covers the following topics:
- Understanding Bytecode
  - How to obtain the bytecode listings
  - How to read the bytecode
  - How the language constructs are mirrored by the compiler: calculation, method calls



# Why Understand Bytecode?

- Understanding the Java bytecode is crucial when debugging and doing performance and memory usage tuning
- Knowing the assembler instructions that are generated from source code can help you know how to code differently to achieve your quality goals
- In addition, the knowledge of bytecode can help you to better understand the Java programming language and know the implementation details for your source code

# What is the Bytecode?

- It is not usual for programmers to program in Java byte code directly, rather they program in Java and then compile their programs into Java byte code
- Each opcode of the bytecode is represented by a single byte

# What is the Bytecode?

- Bytecode is an intermediate representation of a program, halfway between human readable source and machine code
- Bytecode is produced by javac from Java source code files
- Some high-level language features have been removed and do not appear in bytecode
- For example, Java's looping and the if statement are not present, and instead, have become branch instructions in the bytecode

## How to Inspect the Instruction Set

- In order to inspect the output from the Java compiler, we use a program called a disassembler which converts the Java byte code program from a form suitable for efficient interpretation into a textual form suitable for developers to read and study
- Given a compiled Java class in the file `SampleClass.class` the Java disassembler is invoked with a `javap` command such as the following:
  - `javap -c Sampleclass` (The `-c` option disassembles the code)

# JVM Instruction Set

- Java bytecodes are the machine language of the JVM
- In Java source code, the method defines behavior, so after the Classloader loads a class, the stream of bytecodes for each method are stored in the method area of JVM during runtime
- The bytecodes of a method are executed when the method is invoked by the thread during the course of running the application



# JVM Instruction Set

- The stream of bytecodes consists of a one-byte opcode followed by zero or more operands
- Instruction Set format:
  - 1 byte opcode
  - 0 or more bytes of operands
- Each opcode has a mnemonic, for example, `istore_0`.
- JVM instructions are explicitly typed: different codes for instructions for integers, floats, arrays, reference types, etc.

# JVM Instruction Set - Mnemonics

- This is reflected by a naming convention in the first letter of the opcode mnemonics:

- JVM-Type      prefix

- Byte              b
- Short            s
- Integer          i
- Long            l
- Character              c
- Single float          f
- Double float      d
- References          a

- Example: iload   load integer type or fload load float type

# JVM Instruction Set - Mnemonics

- Below is a list of common uses of opcode:
  - Shuffling (pop, swap, dup, ...)
  - Calculating (iadd, isub, imul, idiv, ineg,...)
  - Conversion (d2i, i2b, d2f, i2z,...)
  - Local storage operation (iload, istore,...)
  - Array Operation (arraylength, newarray,...)
  - Object management (get/putfield, invokevirtual, new)
  - Push operation (aconst\_null, iconst\_m1,...)
  - Control flow (nop, goto, jsr, ret, tableswitch,...)
  - Threading (monitorenter, monitorexit,...)

# Loops in Java Source Code

- Consider the following Java Methods :

```
- void testFor() {  
-   for (int i = 0 ; i < 100 ; i++){;}  
- }
```

```
- int i = 0;  
- void testWhile() {  
-   while (i < 100) {  
-       i++;}  
- }
```

Both methods initialize the loop variable i to zero and then increment until i reaches the 100 limit.

- We consider these testFor() and testWhile() methods equivalent in that we can use each method to perform the same task without the need to change the code

## Loops in Java Byte Code

- The Java byte code language has neither a for loop nor a while loop
- Unlike the source code where the loops are visible, in Java byte code the two forms of loops are identical:

Method void testFor()	Method void testWhile()
- 0: iconst_0	0: iconst_0
- 1: istore_1	1: istore_1
- 2: iload_1	2: iload_1
- 3: bipush 100	3: bipush 100
- 5: if_icmpge 14	5: if_icmpge 14
- 8: iinc 1, 1	8: iinc 1, 1
- 11: goto 2	11: goto 2
- 14: return	14: return

# Loops in Java Byte Code

- Below is the byte code generated from the javac compiler for both methods:

- 0: <code>iconst_0</code>	The integer constant zero is pushed on top of the stack
- 1: <code>istore_1</code>	The top of the stack is stored into local variable array one (the variable i)
- 2: <code>iload_1</code>	Load from the local variable
- 3: <code>bipush 100</code>	Load the integer 100 on top of the stack
- 5: <code>if_icmpge 14</code>	Compare the top two items on the stack and jump if (i>=100)
- 8: <code>iinc 1, 1</code>	Increment local variable on by 1 (i++)
- 11: <code>goto 2</code>	
- 14: <code>return</code>	Return void when the end of the method is reached

# Local Storage Operation

- The load and store instructions transfer values between the local variables and the operand stack of a Java Virtual Machine frame
  - iload, iload\_0, fload\_0, ..
  - lstore, lstore, istore\_0..
  - ldc, sipush

# Arithmetic Instructions Mnemonics

- For a Complete list of the Arithmetic instructions, refer to the JVM Specification
- Below is a small subset :
  - Add: iadd, ladd, fadd, dadd
  - Subtract: isub, lsub, fsub, dsub
  - Multiply: imul, lmul, fmul, dmul
  - Divide: idiv, ldiv, fdiv, ddiv



# Java Bytecode Example

- Java Source Code:

```
public class SampleClass {  
    public static int test(){  
        int x=99999;  
        int y=1;  
        int z = x + y;  
        return z;  
    }//end method test  
}//end class SampleClass
```

- # javap -c SampleClass
- Java ByteCode:
  - public static int test();
  - Code:
    - 0: ldc #2 // int 99999
    - 2: istore\_0
    - 3: iconst\_1
    - 4: istore\_1
    - 5: iload\_0
    - 6: iload\_1
    - 7: iadd
    - 8: istore\_2
    - 9: iload\_2
    - 10: ireturn

# Object Management

- Although both class instances and arrays are objects, the Java Virtual Machine creates and manipulates class instances and arrays using distinct sets of instructions:
  - Create a new class instance: `new`
  - Create a new array: `newarray`, `anewarray`, `multianewarray`
- Access fields of classes (static fields, known as class variables) and fields of class instances (non-static fields, known as instance variables):
  - `getstatic`, `putstatic`, `getfield`, `putfield`

## Instruction-Set: Memory Access

- In the JVM, the contents of different “kinds” of memory can be accessed by different kinds of instructions
  - accessing locals and arguments: load and store instructions
  - accessing fields in objects: getfield, putfield
  - accessing static fields: getstatic, putstatic

# Instruction-Set: Memory Access

- Note:

- Static fields are like global variables
- They are allocated in the “method area” where code for methods and representations for classes (including method tables) are also stored
- getfield and putfield access memory in the heap

# The new Operator Example

- Java Source Code:

```
public class SampleClass {  
    public static int testID=100;  
    public static void test()  
    {  
        SampleClass sc=new SampleClass();  
        testID=200;  
    } //end method test  
} //end class SampleClass
```

- In the SampleClass java class, we declare a static variable testID.
- And the the test() method, a new Sampleclass object is instantiated

# The new Operator Example

- Java ByteCode:
- 0: new #2 // class SampleClass
- 3: dup
- 4: invokespecial #3 // Method "<init>:()V"
- 7: astore\_0
- 8: sipush 200
- 11: putstatic #4 // Field testID:I
- 14: return

- In the SampleClass test() method bytecode:

- Line 0: create a new object
- Line 3: duplicate the top operand stack value
- Line 4: invoke the SampleClass constructor
- Line 7: store the newly created object reference into the local variable

# The new Operator Example

- Java ByteCode:
- 0: new #2 // class SampleClass
- 3: dup
- 4: invokespecial #3 // Method "<init>:()V"
- 7: astore\_0
- 8: sipush 200
- 11: putstatic #4 // Field testID:I
- 14: return

- In the SampleClass test() method bytecode:
  - Line 8: push the short 200 into the top of the stack
  - Line 11: set the static field testID to 200
  - Line 14: return the method call

# Instruction-Set

- Method invocation:

- `invokevirtual`: the usual instruction for calling a method on an object. `invokeinterface`: same as `invokevirtual`, but used when the called method is declared in an interface (requires a different kind of method lookup)
- `invokespecial`: for calling things such as constructors, which are not dynamically dispatched (this instruction is also known as `invokenonvirtual`)
- `invokestatic`: for calling methods that have the “static” modifier (these methods are sent to a class, not to an object)

- Returning from methods:

- `return`, `ireturn`, `lreturn`, `areturn`, `freturn`, ...



# Analyze the Bytecode : invokespecial

- Java Source Code:
  - `SampleClass sc=new SampleClass();`
  - The new operator will be translated into 4 instructions by javac compiler
- Java ByteCode:
  - 0: new #2 // class SampleClass
    - 3: dup
    - Duplicate the top operand stack value
    - 4: invokespecial #3 // Method "<init>":()V
    - Call the constructor of the SampleClass
  - Memory for a new instance of SampleClass class is allocated from the garbage-collected heap, and the instance variables of the new object are initialized to their default initial values

# Analyze the Bytecode - Push Short

- Java Source Code:
  - `testID=200;`
  - The statement will be translated into two instructions
- Java ByteCode:
  - `8: sipush 200`
  - Push short value 200 into stack.
  - `11: putstatic #4 //`  
Field testID:I
  - Set static field in class

# Instructions and the “Constant Pool”

- Many JVM instructions have operands which are indexes pointing to an entry in the so-called constant pool
- The constant pool contains all kinds of entries that represent “symbolic” references for “linking”
- This is the way that instructions refer to things such as classes, interfaces, fields, methods, and constants such as string literals and numbers

# Instructions and the “Constant Pool”

- Examples of constant pool entries that exist :
  - Class\_info
  - Fieldref\_info
  - Methodref\_info
  - InterfaceMethodref\_info
  - String
  - Integer
  - Float
  - Long
  - Double
  - Name\_and\_Type\_info
  - Utf8\_info (Unicode characters)

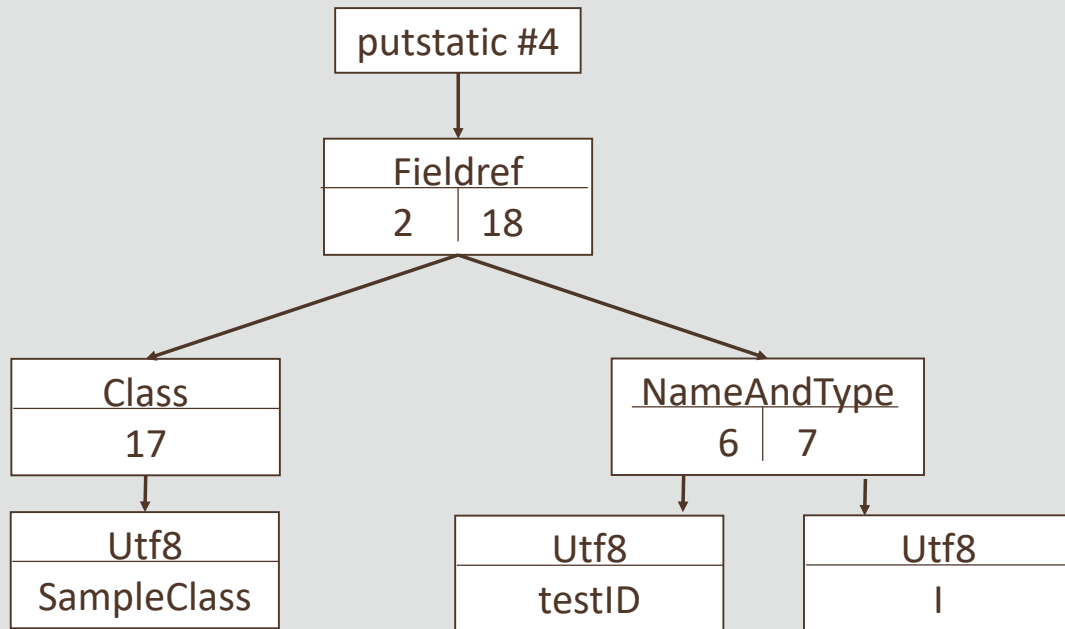
# Analyze the Bytecode

- `#javap -verbose SampleClass // display the constant pool`
- `Java ByteCode: 11: putstatic #4 // Field t`
- `#1 = Methodref #5.#16 // java/lang/Object."<init>":()V`
- `#2 = Class #17 // SampleClass`
- `#3 = Methodref #2.#16 // SampleClass."<init>":()V`
- `#4 = Fieldref #2.#18 // SampleClass.testID:I`
- `#5 = Class #19 // java/lang/Object`
- `#6 = Utf8 testID`
- `#7 = Utf8 I`
- `#8 = Utf8 <init>`
- `#9 = Utf8 ()V`
- `#10 = Utf8 Code`
- `estID:I`

## Analyze the Bytecode – Cont.

- #javap -verbose SampleClass // display the constant pool
- Java ByteCode: 11: putstatic #4 // Field testID:I
- #11 = Utf8 LineNumberTable
- #12 = Utf8 test
- #13 = Utf8 <clinit>
- #14 = Utf8 SourceFile
- #15 = Utf8 SampleClass.java
- #16 = NameAndType #8:#9 // "<init>":()V
- #17 = Utf8 SampleClass
- #18 = NameAndType #6:#7 // testID:I
- #19 = Utf8 java/lang/Object

## Analyze the Bytecode - Cont



# Summary

- In this lesson, you should have learned:
- Understanding Bytecode
  - How to obtain the bytecode listings
  - How to read the bytecode
  - How the language constructs are mirrored by the compiler: calculation, method calls





The logo for Oracle Academy is centered on a light gray background. It features the word "ORACLE" in a bold, orange, sans-serif font. Below it, the word "Academy" is written in a smaller, dark gray, sans-serif font. The entire logo is framed by a thin black border, with dark gray horizontal bars at the top and bottom.

# ORACLE

## Academy